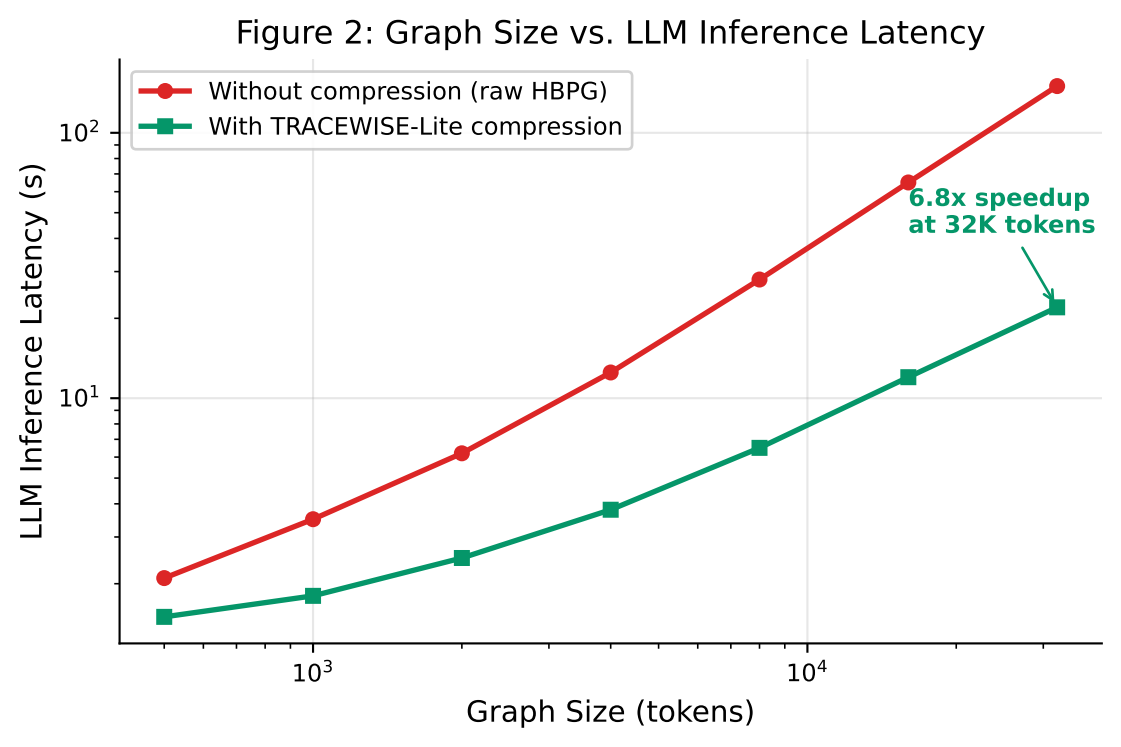
Figure 2: Graph Size vs. LLM Inference Latency
Without compression (raw HBPG)
With TRACEWISE-Lite compression
6.8x speedup
at 32K tokens
LLM Inference Latency (s)
Graph Size (tokens)
10^2
10^1
10^3
10^4